

TTDS: Machine Learning project

Accuracy Improved KNNNeighbour Algorithm

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as plt
from matplotlib import pyplot
import matplotlib.pyplot as plt
```

```
In [2]: data=pd.read_csv("D:/DataSets/diabetes.csv")
```

```
In [3]: data
```

```
Out[3]:
```

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

Dataset Extension

```
In [4]: # Generate synthetic data by doubling the 'label' values
data_synthetic = data.copy()

# Concatenate the original and synthetic DataFrames
df = pd.concat([data, data_synthetic], ignore_index=True)

# Display the extended DataFrame
df
```

Out[4]:

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
1531	10	101	76	48	180	32.9	0.171	63	0
1532	2	122	70	27	0	36.8	0.340	27	0
1533	5	121	72	23	112	26.2	0.245	30	0
1534	1	126	60	0	0	30.1	0.349	47	1
1535	1	93	70	31	0	30.4	0.315	23	0

1536 rows × 9 columns

```
In [5]: df.head(10)
```

Out[5]:

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
5	5	116	74	0	0	25.6	0.201	30	0
6	3	78	50	32	88	31.0	0.248	26	1
7	10	115	0	0	0	35.3	0.134	29	0
8	2	197	70	45	543	30.5	0.158	53	1
9	8	125	96	0	0	0.0	0.232	54	1

```
In [6]: df.tail()
```

Out[6]:

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
1531	10	101	76	48	180	32.9	0.171	63	0
1532	2	122	70	27	0	36.8	0.340	27	0
1533	5	121	72	23	112	26.2	0.245	30	0
1534	1	126	60	0	0	30.1	0.349	47	1
1535	1	93	70	31	0	30.4	0.315	23	0

```
In [7]: df.dtypes
```

```
Out[7]: preg          int64
        glucose       int64
        bp_diastolic   int64
        skin_triceps   int64
        insulin        int64
        bmi            float64
        pedigree       float64
        age            int64
        label          int64
        dtype: object
```

Descriptive Statistics:

```
In [8]: print("Number of Row in the Dataset:", df.shape[0])
        print("Number of Columns in the Dataset:", df.shape[1])
```

```
Number of Row in the Dataset: 1536
Number of Columns in the Dataset: 9
```

```
In [9]: df.head(10)
```

```
Out[9]:
```

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
5	5	116	74	0	0	25.6	0.201	30	0
6	3	78	50	32	88	31.0	0.248	26	1
7	10	115	0	0	0	35.3	0.134	29	0
8	2	197	70	45	543	30.5	0.158	53	1
9	8	125	96	0	0	0.0	0.232	54	1

```
In [10]: df.tail()
```

```
Out[10]:
```

	preg	glucose	bp_diastolic	skin_triceps	insulin	bmi	pedigree	age	label
1531	10	101	76	48	180	32.9	0.171	63	0
1532	2	122	70	27	0	36.8	0.340	27	0
1533	5	121	72	23	112	26.2	0.245	30	0
1534	1	126	60	0	0	30.1	0.349	47	1
1535	1	93	70	31	0	30.4	0.315	23	0

In [11]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1536 entries, 0 to 1535
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   preg             1536 non-null   int64
1   glucose          1536 non-null   int64
2   bp_diastolic     1536 non-null   int64
3   skin_triceps     1536 non-null   int64
4   insulin          1536 non-null   int64
5   bmi              1536 non-null   float64
6   pedigree         1536 non-null   float64
7   age              1536 non-null   int64
8   label            1536 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 108.1 KB
```

In [12]: df.describe().T

Out[12]:

	count	mean	std	min	25%	50%	75%	max
preg	1536.0	3.845052	3.368480	0.000	1.00000	3.0000	6.00000	17.00
glucose	1536.0	120.894531	31.962202	0.000	99.00000	117.0000	140.25000	199.00
bp_diastolic	1536.0	69.105469	19.349501	0.000	62.00000	72.0000	80.00000	122.00
skin_triceps	1536.0	20.536458	15.947021	0.000	0.00000	23.0000	32.00000	99.00
insulin	1536.0	79.799479	115.206457	0.000	0.00000	30.5000	127.25000	846.00
bmi	1536.0	31.992578	7.881592	0.000	27.30000	32.0000	36.60000	67.10
pedigree	1536.0	0.471876	0.331221	0.078	0.24375	0.3725	0.62625	2.42
age	1536.0	33.240885	11.756400	21.000	24.00000	29.0000	41.00000	81.00
label	1536.0	0.348958	0.476796	0.000	0.00000	0.0000	1.00000	1.00

In [13]: *#check label value count*
data.label.value_counts()Out[13]: 0 500
1 268
Name: label, dtype: int64

Missing Values:

In [14]: df.isnull().sum()

Out[14]: preg 0
glucose 0
bp_diastolic 0
skin_triceps 0
insulin 0
bmi 0
pedigree 0
age 0
label 0
dtype: int64

```
In [15]: #check missing Values in the Dataset
missing_data=df.isnull()
for column in missing_data.columns.values.tolist():
    print(column)
    print(missing_data[column].value_counts())
    print("")
```

```
preg
False      1536
Name: preg, dtype: int64
```

```
glucose
False      1536
Name: glucose, dtype: int64
```

```
bp_diastolic
False      1536
Name: bp_diastolic, dtype: int64
```

```
skin_triceps
False      1536
Name: skin_triceps, dtype: int64
```

```
insulin
False      1536
Name: insulin, dtype: int64
```

```
bmi
False      1536
Name: bmi, dtype: int64
```

```
pedigree
False      1536
Name: pedigree, dtype: int64
```


```
age
False      1536
Name: age, dtype: int64
```

```
label
False      1536
Name: label, dtype: int64
```

No missing values found in the dataset, therefore data doesn't need to be drop or replace.

Outliers Analysis

```
In [16]: def show_boxplot(df):  
    plt.rcParams['figure.figsize'] = [14,6]  
    sns.boxplot(data = df, orient="v")  
    plt.title("Outliers Distribution", fontsize = 16)  
    plt.ylabel("Range", fontweight = 'bold')  
    plt.xlabel("Attributes", fontweight = 'bold')  
    show_boxplot(df)
```




```
In [53]: #Function 1st time
def remove_outliers(data):

    df = data.copy()

    for col in list(df.columns):
        Q1 = df[str(col)].quantile(0.25)
        Q3 = df[str(col)].quantile(0.75)
        IQR = Q3 - Q1
        # Define the lower and upper bounds to filter outliers
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        df = df[(df[str(col)] >= lower_bound) & (df[str(col)] <= upper_bound)]


    return df
without_outliers = remove_outliers(df)
show_boxplot(without_outliers)
```



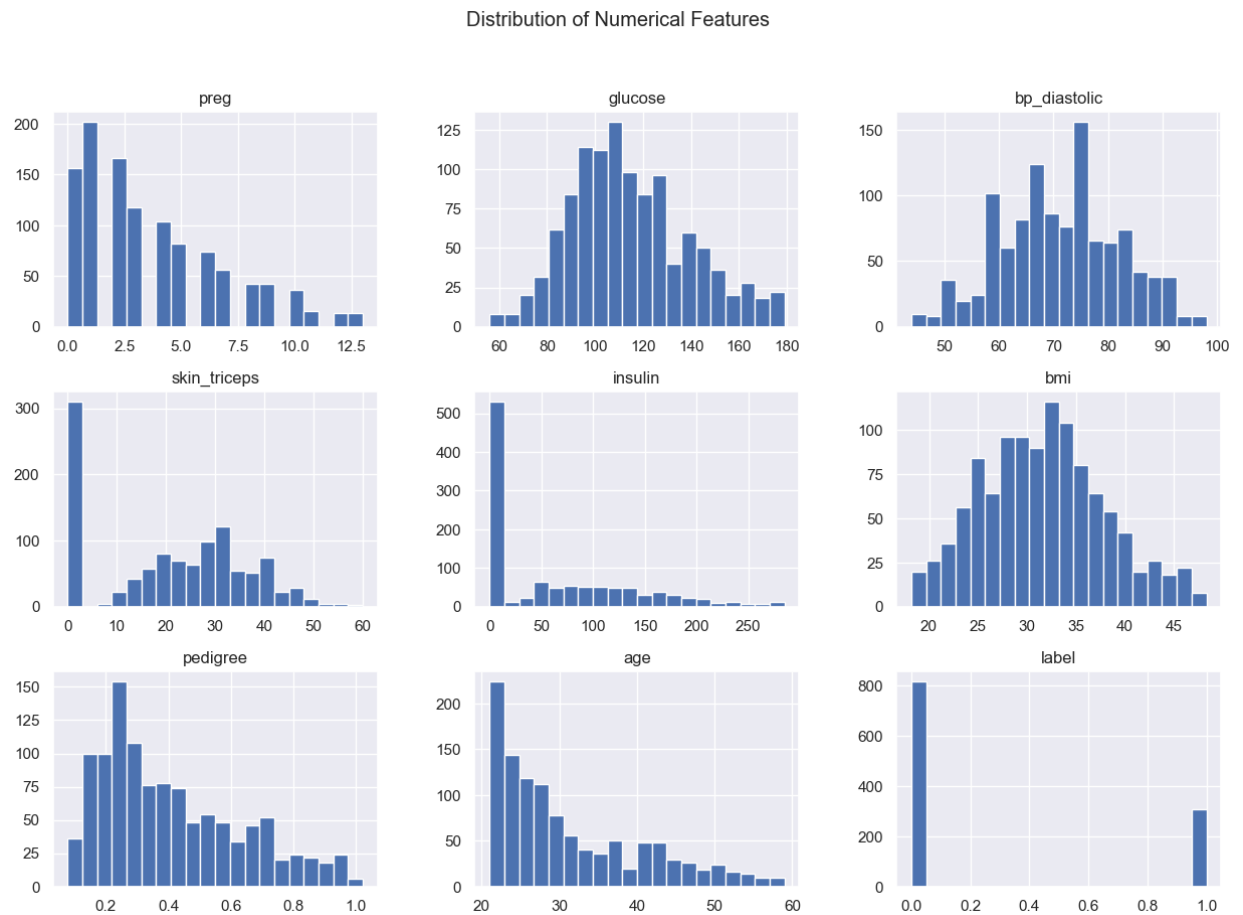
NOTE: Outliers completely removed, after function has been run two to three times.

```
In [54]: df=without_outliers
```

```
In [57]: # Box plots for numerical features
plt.figure(figsize=(15, 10))
df.boxplot()
plt.title('Box Plots for Numerical Features')
plt.show()
```





```
In [55]: # Distribution of numerical features
df.hist(bins=20, figsize=(15, 10))
plt.suptitle('Distribution of Numerical Features')
plt.show()
```




HeatMap

```
In [56]: #get correlations of each features in dataset
corrmat = df.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(6,6))
#plot heat map
sns.heatmap(df[top_corr_features].corr(),annot=True,cmap="RdYlBu", annot_kws={"fontsize": 8});
```



```
# Pair plot
sns.pairplot(df, hue='label')
plt.suptitle('Pair Plot of Features')
plt.show()
```

```
In [58]: # Distribution of categorical features
sns.countplot(x='age', data=df)
plt.title('Distribution of Age')
plt.show()
```



```
In [59]: print(df.shape[0])
print(df.shape[1])
```

```
1122
9
```

```
In [60]: df.label.value_counts()
```

```
Out[60]: 0    816
         1    306
         Name: label, dtype: int64
```

```
In [61]: df.columns
```

```
Out[61]: Index(['preg', 'glucose', 'bp_diastolic', 'skin_triceps', 'insulin', 'bmi',
               'pedigree', 'age', 'label'],
              dtype='object')
```

```
In [62]: cols=list(df.columns)
cols
```

```
Out[62]: ['preg',
          'glucose',
          'bp_diastolic',
          'skin_triceps',
          'insulin',
          'bmi',
          'pedigree',
          'age',
          'label']
```

```
In [63]: df.shape
```

```
Out[63]: (1122, 9)
```

```
In [64]: feature_cols=cols[0:8]
print(feature_cols)
```

```
['preg', 'glucose', 'bp_diastolic', 'skin_triceps', 'insulin', 'bmi', 'pedigree', 'age']
```

```
In [65]: feature_cols=['preg', 'glucose', 'bp_diastolic', 'skin_triceps', 'insulin', 'bmi', 'pedigree',
print(feature_cols)
```

```
['preg', 'glucose', 'bp_diastolic', 'skin_triceps', 'insulin', 'bmi', 'pedigree', 'age']
```

Data Train-Test split

```
In [66]: #Library Call for data split in two portion Train and Test:
from sklearn.model_selection import train_test_split
```

```
In [163]: #dataframe
x=df[feature_cols] #feature
#series
y=df.label
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =0.20, random_state=0)
```

```
In [164]: #Total size of the Training dataset:
print("[XY_Train] dataset Shape:", x_train.shape)

#Total size of the Testing dataset:
print("[XY_Test] dataset Shape:", x_test.shape)
```

```
[XY_Train] dataset Shape: (897, 8)
[XY_Test] dataset Shape: (225, 8)
```

```
In [165]: #get total number of 0 in the actual dataset
count0=df["label"][df.label==0].count()
print("Total Number of 0's in Label:", count0)
```

```
Total Number of 0's in Label: 816
```

```
In [166]: #get total number of 1 in the actual dataset
count1=df["label"][df.label==1].count()
print("Total Number of 1's in Label:", count1)
```

```
Total Number of 1's in Label: 306
```

```
In [167]: #Checking the number of 0's in Training portion of the Dataset:
print("[Y_Train] Total number of [0] in dataset :", len(y_train[y_train==0]))

#Checking the number of 1's in Training portion of the Dataset:
print("[Y_Train] Total number of [1] in dataset :", len(y_train[y_train==1]))
```

```
[Y_Train] Total number of [0] in dataset : 644
[Y_Train] Total number of [1] in dataset : 253
```

```
In [168]: #Checking the number of 0's in Testing portion of the Dataset:
print("[Y_Test] Total number of [0] in dataset :", len(y_test[y_test==0]))

#Checking the number of 1's in Testing portion of the Dataset:
print("[Y_Test] Total number of [1] in dataset :", len(y_test[y_test==1]))
```

```
[Y_Test] Total number of [0] in dataset : 172
[Y_Test] Total number of [1] in dataset : 53
```

```
In [169]: # get total number of 0 in the training dataset
Trcount0 = y_train[y_train==0].count()

# get total number of 1 in the training dataset
Trcount1 = y_train[y_train==1].count()

# Plotting the bar chart
label = ['0', '1']
counts = [Trcount0, Trcount1]

plt.figure(figsize=(4,4))
plt.title('Counts of 0 and 1 in Training Dataset')
plt.bar(label, counts)

# Add annotations to the bars
for i, count in enumerate(counts):
    plt.text(i, count, str(count), ha='center', va='bottom')

plt.show()
```



```
In [170]: # get total number of 0 in the testing dataset
Trcount0 = y_test[y_test==0].count()


# get total number of 1 in the testing dataset
Trcount1 = y_test[y_test==1].count()

# Plotting the bar chart
label = ['0', '1']
counts = [Trcount0, Trcount1]

plt.figure(figsize=(4,4))
plt.title('Counts of 0 and 1 in Testing Dataset')
plt.bar(label, counts)

# Add annotations to the bars
for i, count in enumerate(counts):
    plt.text(i, count, str(count), ha='center', va='bottom')

plt.show()
```



KNeighbour Library call

```
In [171]: from sklearn.neighbors import KNeighborsClassifier

clf= KNeighborsClassifier(n_neighbors=3)
```

```
In [172]: # Train Classifier
model = clf.fit(x_train, y_train)
```

Model

```
In [173]: #Predict the response for test dataset
y_pred = clf.predict(x_test)
```

```
In [174]: y=pd.DataFrame({"Original": y_test, "Predicted": y_pred})  
y.head()
```

Out[174]:

	Original	Predicted
1069	1	0
1271	0	0
548	0	0
1396	0	1
704	0	0

```
In [175]: y.sample(10)
```

Out[175]:

	Original	Predicted
682	0	0
1412	0	0
410	0	0
1470	1	1
1534	1	1
278	0	0
650	0	0
1467	0	0
801	0	0
282	0	0

Confusion Matrics

```
In [176]: # calculate accuracy
from sklearn import metrics


result = metrics.confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(result)

def plt1():
    import seaborn as sns; sns.set()
    plt.figure(figsize=(4,4))
    c_mtx = pd.crosstab(y_test, y_pred, rownames=['Actual'], colnames=['Predicted'])
    sns.heatmap(c_mtx, annot=True, fmt = '.3g')

plt1()
```

Confusion Matrix:

```
[[145  27]
 [ 19  34]]
```



Accuracy Calculation

```
In [177]: #[row, column]
#(Actual, Predict)
TP = result[1, 1]
TN = result[0, 0]
FP = result[0, 1]
FN = result[1, 0]
```



```
In [178]: def EvClsMdl(res):  
    print('Metrics computed from a confusion matrix')  
    print("Accuracy:\t", metrics.accuracy_score(y_test, y_pred))  
    print("Sensitivity:\t", metrics.recall_score(y_test, y_pred))  
    print("Specificity:\t", TN / (TN + FP))  
    print("Precision:\t", metrics.precision_score(y_test, y_pred))  
    print("Classification Error:", 1 - metrics.accuracy_score(y_test, y_pred))  
    print("False_Positive_Rate:", 1 - TN / (TN + FP))  
    print('#####')  
EvClsMdl(result)
```

```
Metrics computed from a confusion matrix  
Accuracy:      0.7955555555555556  
Sensitivity:    0.6415094339622641  
Specificity:    0.8430232558139535  
Precision:      0.5573770491803278  
Classification Error: 0.20444444444444443  
False_Positive_Rate: 0.15697674418604646  
#####
```

In []:

In []: